

ELL 881: Fundamentals of Deep Learning

Lec 03a: Deep Feedforward Networks

Vineet Kumar

August 8, 2018

Table of Contents

Fundamentals

Limitations of Linear Models

Hidden Units

Architecture Design

Table of Contents

Fundamentals

Limitations of Linear Models

Hidden Units

Architecture Design

Deep Feedforward Networks

- ▶ Also known as **Feedforward Neural Networks** or **Multilayer perceptrons** (MLP)
- ▶ In this class, we will refer deep feed forward networks as **MLP**

Deep Feedforward Networks

- ▶ Also known as **Feedforward Neural Networks** or **Multilayer perceptrons** (MLP)
- ▶ In this class, we will refer deep feed forward networks as **MLP**
- ▶ MLP defines mapping $y = f(\mathbf{x}; \boldsymbol{\theta})$

Deep Feedforward Networks

- ▶ Also known as **Feedforward Neural Networks** or **Multilayer perceptrons** (MLP)
- ▶ In this class, we will refer deep feed forward networks as **MLP**
- ▶ MLP defines mapping $y = f(\mathbf{x}; \theta)$
- ▶ f assigns a category y to an input \mathbf{x}

Deep Feedforward Networks

- ▶ Also known as **Feedforward Neural Networks** or **Multilayer perceptrons** (MLP)
- ▶ In this class, we will refer deep feed forward networks as **MLP**
- ▶ MLP defines mapping $y = f(\mathbf{x}; \boldsymbol{\theta})$
- ▶ f assigns a category y to an input \mathbf{x}
- ▶ Goal of MLP is to approximate some true function $f^*(\mathbf{x})$

Deep Feedforward Networks

- ▶ Also known as **Feedforward Neural Networks** or **Multilayer perceptrons** (MLP)
- ▶ In this class, we will refer deep feed forward networks as **MLP**
- ▶ MLP defines mapping $y = f(\mathbf{x}; \boldsymbol{\theta})$
- ▶ f assigns a category y to an input \mathbf{x}
- ▶ Goal of MLP is to approximate some true function $f^*(\mathbf{x})$
- ▶ This is achieved by estimating the parameters $\boldsymbol{\theta}$

Layers as Composition of Functions

- ▶ DFN is represented by composing together many functions

Layers as Composition of Functions

- ▶ DFN is represented by composing together many functions
- ▶ For example, DFN could be constructed by connecting three functions $f^{(1)}$, $f^{(2)}$ and $f^{(3)}$ in a chain such that:

$$f(\mathbf{x}) = f^{(3)}(f^{(2)}(f^{(1)}(\mathbf{x})))$$

Layers as Composition of Functions

- ▶ DFN is represented by composing together many functions
- ▶ For example, DFN could be constructed by connecting three functions $f^{(1)}$, $f^{(2)}$ and $f^{(3)}$ in a chain such that:

$$f(\mathbf{x}) = f^{(3)}(f^{(2)}(f^{(1)}(\mathbf{x})))$$

- ▶ We call $f^{(1)}$ the first layer, $f^{(2)}$ the second layer ...
- ▶ The final layer is called the **output layer**

Multiple Layers

- ▶ MLP can be thought of as a *black box*
- ▶ MLP makes a prediction $f(\mathbf{x})$ for an input \mathbf{x}

Multiple Layers

- ▶ MLP can be thought of as a *black box*
- ▶ MLP makes a prediction $f(\mathbf{x})$ for an input \mathbf{x}
- ▶ MLP is trained by *guiding* the predictions as close as possible to the ground truth category y

Multiple Layers

- ▶ MLP can be thought of as a *black box*
- ▶ MLP makes a prediction $f(\mathbf{x})$ for an input \mathbf{x}
- ▶ MLP is trained by *guiding* the predictions as close as possible to the ground truth category y
- ▶ As we only **observe** the final layer, all intermediate layers are called **hidden**

Multiple Layers

- ▶ MLP can be thought of as a *black box*
- ▶ MLP makes a prediction $f(\mathbf{x})$ for an input \mathbf{x}
- ▶ MLP is trained by *guiding* the predictions as close as possible to the ground truth category y
- ▶ As we only **observe** the final layer, all intermediate layers are called **hidden**
- ▶ Final layer is called the **output layer**




Figure 1: Multi Layer Perceptron
Image: Getting started with Tensorflow, Safari Books, Giancarlo Zaccone

Example code for MLP

Let us see multiple layers in action, via some Tensorflow code!

Notebook: `mlp_example_eager.ipynb`

Table of Contents

Fundamentals

Limitations of Linear Models

Hidden Units

Architecture Design

Limitations of Linear Models: Learning XOR

- ▶ XOR function is an operation on two binary values x_1 and x_2 , which returns 1 only when one of the values is 1.

Limitations of Linear Models: Learning XOR

- ▶ XOR function is an operation on two binary values x_1 and x_2 , which returns 1 only when one of the values is 1.
- ▶ We want to learn XOR function $y = f^*(\mathbf{x})$

Limitations of Linear Models: Learning XOR

- ▶ XOR function is an operation on two binary values x_1 and x_2 , which returns 1 only when one of the values is 1.
- ▶ We want to learn XOR function $y = f^*(\mathbf{x})$
- ▶ Let us try to learn a MLP $y = f(\mathbf{x}; \boldsymbol{\theta})$


Limitations of Linear Models: Learning XOR

- ▶ XOR function is an operation on two binary values x_1 and x_2 , which returns 1 only when one of the values is 1.
- ▶ We want to learn XOR function $y = f^*(\mathbf{x})$
- ▶ Let us try to learn a MLP $y = f(\mathbf{x}; \boldsymbol{\theta})$
- ▶ We only care about $X = \{[0, 0]^T, [0, 1]^T, [1, 0]^T, [1, 1]^T\}$

Code: Linear XOR Model

Notebook: xor_eager_keras.ipynb

Why does a linear XOR Model fail?



- ▶ We cannot find a line which separates $label = 1$ from $label = 0$

Figure 2: The data points are not separable via a linear function

Image Courtesy:

<https://medium.com/@jayeshbahire/the-xor-problem-in-neural-networks-50006411840b>

Why does a linear XOR Model fail?




Figure 2: The data points are not separable via a linear function

Image Courtesy:

<https://medium.com/@jayeshbahire/the-xor-problem-in-neural-networks-50006411840b>

- ▶ We cannot find a line which separates $label = 1$ from $label = 0$
- ▶ Thus, the solution is to add a **non linear** layer

XOR Model: Adding multiple layers

- ▶ We define a two layer network: one *hidden layer* and one *output* layer

XOR Model: Adding multiple layers

- ▶ We define a two layer network: one *hidden layer* and one *output layer*
- ▶ First layer $\mathbf{h} = f^{(1)}(\mathbf{x}; W, c)$
- ▶ Second layer $y = f^{(2)}(\mathbf{h}; w, b)$

XOR Model: Adding multiple layers

- ▶ We define a two layer network: one *hidden layer* and one *output layer*
- ▶ First layer $\mathbf{h} = f^{(1)}(\mathbf{x}; W, c)$
- ▶ Second layer $y = f^{(2)}(\mathbf{h}; w, b)$
- ▶ If both $f^{(1)}$ and $f^{(2)}$ are linear, we effectively have only **one linear layer!** Why?

XOR Model: Adding multiple layers

- ▶ We define a two layer network: one *hidden layer* and one *output layer*
- ▶ First layer $\mathbf{h} = f^{(1)}(\mathbf{x}; W, c)$
- ▶ Second layer $y = f^{(2)}(\mathbf{h}; w, b)$
- ▶ If both $f^{(1)}$ and $f^{(2)}$ are linear, we effectively have only **one linear layer!** Why?
- ▶ $f^{(1)}(x) = W^T \mathbf{x}$; $f^{(2)}(h) = h^T \mathbf{w}$; Thus, $f(\mathbf{x}) = \mathbf{w}^T W^T \mathbf{x}$

XOR Model: Adding multiple layers

Notebook: xor_eager_keras_ml.ipynb

XOR Model: Adding multiple layers

Notebook: xor_eager_keras_ml.ipynb

Add Non-linear Layer

XOR Model: Adding Non-Linear Layer

- ▶ We apply a nonlinear function g , such that

$$\mathbf{h} = g(W^T \mathbf{x} + \mathbf{c})$$

- ▶ Note, that g is applied elementwise

XOR Model: Adding Non-Linear Layer

- ▶ We apply a nonlinear function g , such that

$$\mathbf{h} = g(W^T \mathbf{x} + \mathbf{c})$$

- ▶ Note, that g is applied elementwise
- ▶ The default recommendation is to use Rectified Linear Unit **ReLU**.
- ▶ $g(z) = \max\{0, z\}$

XOR Model: Adding Non-Linear Layer

- ▶ We apply a nonlinear function g , such that

$$\mathbf{h} = g(W^T \mathbf{x} + \mathbf{c})$$

- ▶ Note, that g is applied elementwise
- ▶ The default recommendation is to use Rectified Linear Unit **ReLU**.
- ▶ $g(z) = \max\{0, z\}$
- ▶ Note that ReLU is a piecewise linear function, and still has easy to compute derivatives.
- ▶ We will talk more about ReLU in next section.




Figure 3: ReLU

Image Courtesy:

http://www.deeplearningbook.org/slides/06_mlp.pdf

XOR Model: Adding Non-Linear Layer

Notebook: `xor_eager_keras_ml_relu.ipynb`

Table of Contents

Fundamentals

Limitations of Linear Models

Hidden Units

Architecture Design

Non differentiable Hidden Unit

- ▶ Design of hidden units is still an **active research area**

Non differentiable Hidden Unit

- ▶ Design of hidden units is still an **active research area**
- ▶ ReLU is an excellent default choice

Non differentiable Hidden Unit

- ▶ Design of hidden units is still an **active research area**
- ▶ ReLU is an excellent default choice
- ▶ ReLU $g(z) = \max\{0, z\}$ is **not differentiable** at all points.
How can we still use it for learning?

Non differentiable Hidden Unit

- ▶ Design of hidden units is still an **active research area**
- ▶ ReLU is an excellent default choice
- ▶ ReLU $g(z) = \max\{0, z\}$ is **not differentiable** at all points.
How can we still use it for learning?
- ▶ Neural Network training does not usually reach a local minimum, and it is okay to not have a gradient defined at 0.

Non differentiable Hidden Unit

- ▶ Design of hidden units is still an **active research area**
- ▶ ReLU is an excellent default choice
- ▶ ReLU $g(z) = \max\{0, z\}$ is **not differentiable** at all points.
How can we still use it for learning?
- ▶ Neural Network training does not usually reach a local minimum, and it is okay to not have a gradient defined at 0.
- ▶ Key point to remember: ReLU is not differentiable at 0 but it can be used as its left and right derivative are defined.

Hidden Unit

- ▶ Most hidden units, first do an affine transformation of the input

$$z = W^T \mathbf{x} + \mathbf{b}$$

Hidden Unit

- ▶ Most hidden units, first do an affine transformation of the input

$$z = W^T \mathbf{x} + \mathbf{b}$$

- ▶ They later apply an element wise *nonlinear* function $g(z)$ such as ReLU

Hidden Unit

- ▶ Most hidden units, first do an affine transformation of the input

$$z = W^T \mathbf{x} + \mathbf{b}$$

- ▶ They later apply an element wise *nonlinear* function $g(z)$ such as ReLU
- ▶ Hidden units usually only differ in choice of the function g


More on ReLU

- ▶ ReLU is similar to a linear unit
- ▶ Gradients are large and consistent even for small values of input!
- ▶ This makes learning easier

More on ReLU

- ▶ ReLU is similar to a linear unit
- ▶ Gradients are large and consistent even for small values of input!
- ▶ This makes learning easier
- ▶ It is important to **initialize the biases with small constant values** such as 0.1
- ▶ This allows ReLU units activate initially, and allow them to pass gradients through!

Logistic Sigmoid and Hyperbolic Tangent



- Unlike ReLU, sigmoid suffers from *saturation*

Figure 4: Sigmoid

Image Courtesy:

https://en.wikipedia.org/wiki/Logistic_function




Figure 5: Tanh

Image Courtesy:

<http://mathworld.wolfram.com/HyperbolicTangent.html>

Logistic Sigmoid and Hyperbolic Tangent




Figure 4: Sigmoid

Image Courtesy:

https://en.wikipedia.org/wiki/Logistic_function




Figure 5: Tanh

Image Courtesy:

<http://mathworld.wolfram.com/HyperbolicTangent.html>

- ▶ Unlike ReLU, sigmoid suffers from *saturation*
- ▶ When z is very positive σ saturates to a high value
- ▶ When z is very negative σ saturates to a low value
- ▶ Sigmoid is only strongly sensitive to the input near zero.
- ▶ Thus, use of sigmoids for MLP is **discouraged**
- ▶ A better alternative is to use *tanh*, which behaves like a linear function, when activations are small.

Table of Contents

Fundamentals

Limitations of Linear Models

Hidden Units

Architecture Design

Architecture Design Choices




Figure 6: Architecture Design Choices

Image Courtesy: http://www.deeplearningbook.org/slides/06_mlp.pdf

Universal Approximation Properties & Depth

- ▶ **Universal Approximation Theorem** states that one hidden layer (such as ReLU or sigmoid) is enough to represent an approximation of any function

Universal Approximation Properties & Depth

- ▶ **Universal Approximation Theorem** states that one hidden layer (such as ReLU or sigmoid) is enough to represent an approximation of any function
- ▶ However, this theorem only talks about **representing** a function, and not **learning** it!

Universal Approximation Properties & Depth

- ▶ **Universal Approximation Theorem** states that one hidden layer (such as ReLU or sigmoid) is enough to represent an approximation of any function
- ▶ However, this theorem only talks about **representing** a function, and not **learning** it!
- ▶ MLP only provide a guarantee that there exists some MLP which can estimate a function

Universal Approximation Properties & Depth

- ▶ **Universal Approximation Theorem** states that one hidden layer (such as ReLU or sigmoid) is enough to represent an approximation of any function
- ▶ However, this theorem only talks about **representing** a function, and not **learning** it!
- ▶ MLP only provide a guarantee that there exists some MLP which can estimate a function
- ▶ In practise, using deeper models can reduce the number of units required to learn a function.

Universal Approximation Properties & Depth

- ▶ **Universal Approximation Theorem** states that one hidden layer (such as ReLU or sigmoid) is enough to represent an approximation of any function
- ▶ However, this theorem only talks about **representing** a function, and not **learning** it!
- ▶ MLP only provide a guarantee that there exists some MLP which can estimate a function
- ▶ In practise, using deeper models can reduce the number of units required to learn a function.
- ▶ Another reason to select deeper network is to define a model in terms of composition of simpler functions.